

Task1:

A = load '/hadoopdata/pig/test.txt/' using PigStorage as (line:chararray);

grunt> B = FOREACH A GENERATE flatten(TOKENIZE((chararray)\$0)) as word;

grunt> dump B;

Output of B

```
()
(Hi)
(this)
(is)
(big)
(data)
(class)
(Hi)
(this)
(is)
(big)
(data)
(class)
()
(Hi)
(this)
(is)
(big)
(data)
(class)
(Hi)
(this)
(big)
(data)
(class)
(Hi)
(this)
(big)
(data)
(class)
(Hi)
(this)
(big)
(data)
(class)
grunt> You have new mail in /var/spool/mail/acadgild
```

C = group B by word;

Dump C;

```
2018-08-02 10:31:04,880 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
2018-08-02 10:31:04,885 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code.
2018-08-02 10:31:04,912 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2018-08-02 10:31:04,912 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(Hi, {(Hi), (Hi), (Hi), (Hi)})
(is, {(is), (is), (is)})
(big, {(big), (big), (big), (big), (big)})
(data, {(data), (data), (data), (data), (data)})
(this, {(this), (this), (this), (this), (this)})
(class, {(class), (class), (class), (class), (class)})
(), {( ), ()}
grunt>
```

D = foreach C generate group, COUNT(B);

```
2018-08-02 10:35:07,251 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(Hi, 5)
(is, 3)
(big, 5)
(data, 5)
(this, 5)
(class, 5)
(), 0)
grunt>
```

Task2: We have employee_details and employee_expenses files. Use local mode while running Pig and write Pig Latin script to get below results:

employee_details (EmpID,Name,Salary,EmployeeRating)

```
emp = Load '/hadoopdata/pig/employee_details.txt/' using PigStorage(',') AS (emp_id :int,  
emp_name:chararray,emp_salary:int, emp_rating:int);
```

```
dump emp;
```

```
2016-08-02 11:22:50.105 [main] INFO org.apache.pig.backend.hadoop.ExecutedOnEngine:Total input paths to process: 1  
(101,Amitabh,20000,1)  
(102,Shahrukh,10000,2)  
(103,Akshay,11000,3)  
(104,Anubhav,5000,4)  
(105,Pawan,2500,5)  
(106,Aamir,25000,1)  
(107,Salman,17500,2)  
(108,Ranbir,14000,3)  
(109,Katrina,1000,4)  
(110,Priyanka,2000,5)  
(111,Tushar,500,1)  
(112,Ajay,5000,2)  
(113,Jubeen,1000,1)  
(114,Madhuri,2000,2)  
grunt>
```

employee_expenses(EmpID,Expense)

```
emp_expenses = LOAD '/hadoopdata/pig/employee_expense.txt/' AS (emp_id: int, expenses: int);
```

```
2016-08-02 11:20:22.105 [main] INFO org.apache.pig.backend.hadoop.ExecutedOnEngine:Total input paths to process: 1  
(101,200)  
(102,100)  
(110,400)  
(114,200)  
(119,200)  
(105,100)  
(101,100)  
(104,300)  
(102,400)  
grunt>
```

(a) Top 5 employees (employee id and employee name) with highest rating. (In case two employees have same rating, employee with name coming first in dictionary should get preference)

```
emp_rating1 = order emp by emp_rating DESC, emp_name ASC;
```

```
(105,Pawan,2500,5)  
(110,Priyanka,2000,5)  
(104,Anubhav,5000,4)  
(109,Katrina,1000,4)  
(103,Akshay,11000,3)  
(108,Ranbir,14000,3)  
(112,Ajay,5000,2)  
(114,Madhuri,2000,2)  
(107,Salman,17500,2)  
(102,Shahrukh,10000,2)  
(106,Aamir,25000,1)  
(101,Amitabh,20000,1)  
(113,Jubeen,1000,1)  
(111,Tushar,500,1)
```

```
emp_rating_limit = LIMIT emp_rating1 5;
```

```
Dump emp_rating_limit;
```

```
(105,Pawan,2500,5)
(110,Priyanka,2000,5)
(104,Anubhav,5000,4)
(109,Katrina,1000,4)
(103,Akshay,11000,3)
grunt>
```

Top 3 employees (employee id and employee name) with highest salary, whose employee id is an odd number. (In case two employees have same salary, employee with name coming first in dictionary should get preference)

```
emp_sal = ORDER emp by emp_salary DESC, emp_name ASC;
```

```
(106,Aamir,25000,1)
(101,Amitabh,20000,1)
(107,Salman,17500,2)
(108,Ranbir,14000,3)
(103,Akshay,11000,3)
(102,Shahrukh,10000,2)
(112,Ajay,5000,2)
(104,Anubhav,5000,4)
(105,Pawan,2500,5)
(114,Madhuri,2000,2)
(110,Priyanka,2000,5)
(113,Jubeen,1000,1)
(109,Katrina,1000,4)
(111,Tushar,500,1)
grunt> emp_sal_id = FILTER emp_sal by emp_id%2==1;
grunt> dump emp_sal_id;
```

```
emp_sal_id = FILTER emp_sal by emp_id%2==1;
```

```
dump emp_sal_id;
```

```
2018-08-02 17:21:20.448 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(101,Amitabh,20000,1)
(107,Salman,17500,2)
(103,Akshay,11000,3)
(105,Pawan,2500,5)
(113,Jubeen,1000,1)
(109,Katrina,1000,4)
(111,Tushar,500,1)
grunt>
```

```
emp_final_limit = LIMIT emp_sal_id 3;
```

```
grunt> dump emp_final_id ;
```

```
(101,Amitabh,20000,1)
(107,Salman,17500,2)
(103,Akshay,11000,3)
grunt>
```

(c) Employee (employee id and employee name) with maximum expense (In case two employees have same expense, employee with name coming first in dictionary should get preference)

```
emp = Load '/hadoopdata/pig/employee_details.txt/' using PigStorage(',') AS (emp_id :int,
emp_name:chararray,emp_salary:int, emp_rating:int);

empexpenses = LOAD '/hadoopdata/pig/employee_expense.txt/' AS (emp_id :int, expense_emp
:int);

emp_exp = JOIN emp BY emp_id , empexpenses BY emp_id;

grunt> maxexpense = ORDER emp_exp BY empexpenses::expense_emp DESC;

grunt> dump maxexpense ;
```

```
2018-08-02 18:10:03,372 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(101,Amitabh,20000,1,101,100)
(101,Amitabh,20000,1,101,200)
(102,Shahrukh,10000,2,102,400)
(102,Shahrukh,10000,2,102,100)
(104,Anubhav,5000,4,104,300)
(105,Pawan,2500,5,105,100)
(110,Priyanka,2000,5,110,400)
(114,Madhuri,2000,2,114,200)
grunt> emp_exp_max = JOIN emp BY emp_id , emp_expenses_max BY emp_id; You have new mail in /var/spool/mail/acadgild
```

```
Limitmaxexpense = LIMIT maxexpense 1;
```

```
2018-08-02 19:35:38,410 [main] INFO
(110,Priyanka,2000,5,110,400)
grunt>
```

```
Limitmaxexpensefinal= FOREACH Limitmaxexpense generate emp:emp_id, emp:emp_name;
```

```
Dump Limitmaxexpensefinal;
```

```
2018-08-02 20:02:39,990 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
2018-08-02 20:02:40,004 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code.
2018-08-02 20:02:40,032 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2018-08-02 20:02:40,032 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(110,Priyanka)
grunt>
```

d) (d) List of employees (employee id and employee name) having entries in employee_expenses file.

```
grunt> A = Load '/hadoopdata/pig/employee_details.txt/' using PigStorage(',') AS (emp_id :int,
emp_name:chararray,emp_salary:int, emp_rating:int);
```

```
B = LOAD '/hadoopdata/pig/employee_expense.txt/' AS (emp_id :int, expense_emp :int);
```

```
emp_expenses = JOIN A BY emp_id, B BY emp_id;
```

```
emp_expense_data = FOREACH emp_expenses GENERATE emp::emp_id,emp:: emp_name;
```

```
grunt> emp_expenses = JOIN A BY emp_id, B BY emp_id;
grunt> emp_expense_data = FOREACH emp_expenses GENERATE A::emp_id, A::emp_name;
grunt> emp_exp_distinct_data = DISTINCT emp_expense_data;
grunt> dump emp_exp_distinct_data;
```

```
2018-08-02 20:55:05/501 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce/libinputformat - Total input paths to process : 1
(101,Amitabh)
(102,Shahrukh)
(104,Anubhav)
(105,Pawan)
(110,Priyanka)
(114,Madhuri)
grunt>
```

(e) List of employees (employee id and employee name) having no entry in employee_expenses file.

A = Load '/hadoopdata/pig/employee_details.txt/' using PigStorage(',') AS (emp_id :int, emp_name:chararray,emp_salary:int, emp_rating:int);

B = LOAD '/hadoopdata/pig/employee_expense.txt/' AS (emp_id :int, expense_emp :int);

```
emp_without_exp = JOIN A BY emp_id LEFT OUTER, B BY emp_id;
```

```
emp_without_exp_filter = FILTER emp_without_exp BY B::emp_id is null;
```

```
emp_without_exp_filter_data = FOREACH emp_without_exp_filter GENERATE A::emp_id, A::emp_name;
```

```
dump emp_without_exp_filter_data;
```

```
2018-08-02 21:04:51/507 [main] INFO org.apache.hadoop.mapreduce.libinputformat - Total input paths to process : 1
2018-08-02 21:04:51,987 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(103,Akshay)
(106,Aamir)
(107,Salman)
(108,Ranbir)
(109,Katrina)
(111,Tushar)
(112,Ajay)
(113,Jubeen)
grunt>
```

Task 3: Implementing the use case: aviation-data-analysis-using-apache-pig.

```
grunt> Register '/home/master/Downloads/piggybank.jar';
2018-02-13 23:35:52,828 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead,
use fs.defaultFS
grunt> █
```

Problem1: Top 5 most visited destinations.

```
grunt> A = load '/user/pig/DelayedFlights.csv' USING org.apache.pig.piggybank.storage.CSVExcelStorage('','NO_MULTILINE','UNIX','SKIP_INPUT_HEADER');
2018-02-13 23:37:08,751 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
grunt>
grunt> B = foreach A generate (int)$1 as year, (int)$10 as flight_num, (chararray)$17 as origin, (chararray) $18 as dest;
grunt>
grunt> C = filter B by dest is not null;
grunt>
grunt> D = group C by dest;
grunt>
grunt> E = foreach D generate group, COUNT(C.dest);
grunt>
grunt> F = order E by $1 DESC;
grunt>
grunt> Result = LIMIT F 5;
grunt>
grunt> A1 = load '/user/pig/airports.csv' USING org.apache.pig.piggybank.storage.CSVExcelStorage('','NO_MULTILINE','UNIX','SKIP_INPUT_HEADER');
2018-02-13 23:37:13,672 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
grunt>
grunt> A2 = foreach A1 generate (chararray)$0 as dest, (chararray)$2 as city, (chararray)$4 as country;
grunt>
grunt> joined_table = join Result by $0, A2 by dest;
grunt>
grunt> dump joined_table; █
```

```

2018-02-13 23:43:54,029 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2018-02-13 23:43:54,029 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(ATL,106898,ATL,Atlanta,USA)
(DEN,63003,DEN,Denver,USA)
(DFW,70657,DFW,Dallas-Fort Worth,USA)
(LAX,59969,LAX,Los Angeles,USA)
(ORD,108984,ORD,Chicago,USA)
grunt>
grunt>
grunt>
grunt>
grunt>

```

Problem2: Which month has seen the most number of cancellations due to bad weather.

```

grunt> A = load '/user/pig/DelayedFlights.csv' USING org.apache.pig.piggybank.storage.CSVExcelStorage(',', 'NO_MULTILINE', 'UNIX', 'SKIP_INPUT_HEADER');
2018-02-13 23:48:22,783 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
grunt>
grunt> B = foreach A generate (int)$2 as month, (int)$10 as flight_num, (int)$22 as cancelled, (chararray)$23 as cancel_code;
grunt>
grunt> C = filter B by cancelled == 1 AND cancel_code == 'B';
grunt>
grunt> D = group C by month;
grunt>
grunt> E = foreach D generate group, COUNT(C.cancelled);
grunt>
grunt> F = order E by $1 DESC;
grunt>
grunt> Result = limit F 1;
grunt>
grunt> dump Result;

```

```

2018-02-13 23:43:54,029 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2018-02-13 23:43:54,029 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(ATL,106898,ATL,Atlanta,USA)
(DEN,63003,DEN,Denver,USA)
(DFW,70657,DFW,Dallas-Fort Worth,USA)
(LAX,59969,LAX,Los Angeles,USA)
(ORD,108984,ORD,Chicago,USA)
grunt>
grunt>
grunt>
grunt>
grunt>

```

Problem2: Which month has seen the most number of cancellations due to bad weather.

```
grunt> A = load '/user/pig/DelayedFlights.csv' USING org.apache.pig.piggybank.storage.CSVExcelStorage('','NO_MULTILINE','UNIX','SKIP_INPUT_HEADER');
2018-02-13 23:48:22,783 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
grunt>
grunt> B = foreach A generate (int)$2 as month,(int)$10 as flight_num,(int)$22 as cancelled,(chararray)$23 as cancel_code;
grunt>
grunt> C = filter B by cancelled == 1 AND cancel_code == 'B';
grunt>
grunt> D = group C by month;
grunt>
grunt> E = foreach D generate group, COUNT(C.cancelled);
grunt>
grunt> F = order E by $1 DESC;
grunt>
grunt> Result = limit F 1;
grunt>
grunt> dump Result;
```

```
Job Stats (time in seconds):
JobId  Maps  Reduces MaxMapTime  MinMapTime  AvgMapTime  MedianMapTime  MaxReduceTime  MinReduceTime  AvgReduceTime  MedianReduceTime
eature Outputs
job_local150561383_0007 1 1 n/a n/a n/a n/a n/a n/a F SAMPLER
job_local1516058916_0009 1 1 n/a n/a n/a n/a n/a n/a F hdfs://localhost:54310/tmp2063473579,
/tmp2063473579,
job_local1590884819_0006 2 1 n/a n/a n/a n/a n/a n/a n/a A,B,C,D,E GROUP_BY,COMBINER
job_local1999420129_0008 1 1 n/a n/a n/a n/a n/a n/a F ORDER_BY,COMBINER

Input(s):
Successfully read 1936758 records (1475388323 bytes) from: "/user/pig/DelayedFlights.csv"
```

```
Job Stats (time in seconds):
JobId  Maps  Reduces MaxMapTime  MinMapTime  AvgMapTime  MedianMapTime  MaxReduceTime  MinReduceTime  AvgReduceTime  MedianReduceTime
eature Outputs
job_local150561383_0007 1 1 n/a n/a n/a n/a n/a n/a F SAMPLER
job_local1516058916_0009 1 1 n/a n/a n/a n/a n/a n/a F hdfs://localhost:54310/tmp2063473579,
/tmp2063473579,
job_local1590884819_0006 2 1 n/a n/a n/a n/a n/a n/a n/a A,B,C,D,E GROUP_BY,COMBINER
job_local1999420129_0008 1 1 n/a n/a n/a n/a n/a n/a F ORDER_BY,COMBINER

Input(s):
Successfully read 1936758 records (1475388323 bytes) from: "/user/pig/DelayedFlights.csv"

Output(s):
Successfully stored 1 records (100671595 bytes) in: "hdfs://localhost:54310/tmp/temp201885068/tmp2063473579"

Counters:
Total records written : 1
Total bytes written : 100671595
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0
```

OUTPUT:

Problem3: Top ten origins with the highest AVG departure delay.

```
grunt> A = load '/user/pig/DelayedFlights.csv' USING org.apache.pig.piggybank.storage.CSVExcelStorage('','NO_MULTILINE','UNIX','SKIP_INPUT_HEADER');
2018-02-13 23:53:53,710 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
grunt>
grunt> B1 = foreach A generate (int)$16 as dep_delay, (chararray)$17 as origin;
grunt>
grunt> C1 = filter B1 by (dep_delay is not null) AND (origin is not null);
grunt>
grunt> D1 = group C1 by origin;
grunt>
grunt> E1 = foreach D1 generate group, AVG(C1.dep_delay);
grunt>
grunt> Result = order E1 by $1 DESC;
grunt>
grunt> Top_ten = limit Result 10;
grunt>
grunt>
grunt> Lookup = load '/user/pig/airports.csv' USING org.apache.pig.piggybank.storage.CSVExcelStorage('','NO_MULTILINE','UNIX','SKIP_INPUT_HEADER');
2018-02-13 23:53:58,798 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
grunt>
grunt> Lookup1 = foreach Lookup generate (chararray)$0 as origin, (chararray)$2 as city, (chararray)$4 as country;
grunt>
grunt> Joined = join Lookup1 by origin, Top_ten by $0;
grunt>
grunt> Final = foreach Joined generate $0,$1,$2,$4;
grunt>
grunt> Final_Result = ORDER Final by $3 DESC;
grunt>
grunt> dump Final_Result;
```



```

Success!

Job Stats (time in seconds):
JobId Maps Reduces MaxMapTime MinMapTime AvgMapTime MedianMapTime MaxReduceTime MinReduceTime AvgReduceTime MedianReduceTime
ature Outputs
job_local1098027478_0015 1 1 n/a n/a n/a n/a n/a n/a Final Result SAMPLER
job_local1106705881_0012 1 1 n/a n/a n/a n/a n/a n/a Result ORDER_BY,COMBINER
job_local11263308022_0016 1 1 n/a n/a n/a n/a n/a n/a Final Result ORDER_BY hdfs://loc
mp/temp201885068/tmp650864016,
job_local1344424114_0014 2 1 n/a n/a n/a n/a n/a n/a Final,Joined,Lookup,Lookup1 HASH_JOIN
job_local1912291826_0010 2 1 n/a n/a n/a n/a n/a n/a A,B1,C1,D1,E1 GROUP_BY,COMBINER
job_local191200221_0011 1 1 n/a n/a n/a n/a n/a n/a Result SAMPLER
job_local1969371809_0013 1 1 n/a n/a n/a n/a n/a n/a Result

Input(s):
Successfully read 1936758 records (2286399089 bytes) from: "/user/pig/DelayedFlights.csv"
Successfully read 3376 records from: "/user/pig/airports.csv"

Output(s):
Successfully stored 10 records (178978395 bytes) in: "hdfs://localhost:54310/tmp/temp201885068/tmp650864016"

Counters:
Total records written : 10
Total bytes written : 178978395
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:

```

Output:-

```

2018-02-13 23:55:26,551 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
2018-02-13 23:55:26,551 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2018-02-13 23:55:26,552 [main] WARN org.apache.pig.data.SchemaTupleBackend - SchemaTupleBackend has already been initialized
2018-02-13 23:55:26,555 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2018-02-13 23:55:26,556 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(CHX,Hancock,USA,116.1470588235294)
(PLN,Pellston,USA,93.76190476190476)
(SPI,Springfield,USA,83.84873949579831)
(ALO,Waterloo,USA,82.2258064516129)
(MQT,NA,USA,79.55665024630542)
(ACY,Atlantic City,USA,79.3103448275862)
(MOT,Minot,USA,78.66165413533835)
(HHH,NA,USA,76.53005464480874)
(EGE,Eagle,USA,74.12891986062718)
(BGM,Binghamton,USA,73.15533980582525)

```

Problem4: Which route (origin & destination) has seen the maximum diversion.

```

grunt> A = load '/user/pig/DelayedFlights.csv' USING org.apache.pig.piggybank.storage.CSVExcelStorage(',',',NO_MULTILINE','UNIK','SKIP_INPUT_HEADER');
2018-02-13 23:58:22,279 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
grunt>
grunt> B = FOREACH A GENERATE (chararray)$17 as origin, (chararray)$18 as dest, (int)$24 as diversion;
grunt>
grunt> C = FILTER B BY (origin is not null) AND (dest is not null) AND (diversion == 1);
grunt>
grunt> D = GROUP C by (origin,dest);
grunt>
grunt> E = FOREACH D generate group, COUNT(C.diversion);
grunt>
grunt> F = ORDER E BY $1 DESC;
grunt>
grunt> Result = limit F 10;
grunt>
grunt> dump Result;

```

```

Success!

Job Stats (time in seconds):
JobId  Maps    Reduces MaxMapTime    MinMapTime    AvgMapTime    MedianMapTime    MaxReduceTime    MinReduceTime    AvgReduceTime    MedianReduceTime
eature Outputs
job_local1143279686_0020    1      1      n/a      n/a      n/a      n/a      n/a      n/a      F      hdfs://localhost:54310/tmp/temp-1037999737,
/tmp-1037999737,
job_local1343752920_0017    2      1      n/a      n/a      n/a      n/a      n/a      n/a      A,B,C,D,E    GROUP_BY_COMBINER
job_local1862513453_0018    1      1      n/a      n/a      n/a      n/a      n/a      n/a      F      SAMPLER
job_local2147353027_0019    1      1      n/a      n/a      n/a      n/a      n/a      n/a      F      ORDER_BY_COMBINER

Input(s):
Successfully read 1936758 records (3148422413 bytes) from: "/user/pig/DelayedFlights.csv"

Output(s):
Successfully stored 10 records (223816222 bytes) in: "hdfs://localhost:54310/tmp/temp201885068/tmp-1037999737"

Counters:
Total records written : 10
Total bytes written : 223816222
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

```

Output:-

```

2018-02-14 00:01:30,655 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
2018-02-14 00:01:30,657 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2018-02-14 00:01:30,658 [main] WARN org.apache.pig.data.SchemaTupleBackend - SchemaTupleBackend has already been initialized
2018-02-14 00:01:30,673 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2018-02-14 00:01:30,674 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
((ORD,LGA),39)
((DAL,HOU),35)
((DFW,LGA),33)
((ATL,LGA),32)
((ORD,SNA),31)
((SLC,SUN),31)
((MIA,LGA),31)
((BUR,JFK),29)
((HRL,HOU),28)
((BUR,DFW),25)

```

```

2018-02-13 23:50:17,897 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
2018-02-13 23:50:17,898 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2018-02-13 23:50:17,899 [main] WARN org.apache.pig.data.SchemaTupleBackend - SchemaTupleBackend has already been initialized
2018-02-13 23:50:17,931 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2018-02-13 23:50:17,932 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(12,250)

```